

Software

ConvNets and Transfer Learning

Legal Notices and Disclaimers

This presentation is for informational purposes only. INTEL MAKES NO WARRANTIES, EXPRESS OR IMPLIED, IN THIS SUMMARY.

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. Check with your system manufacturer or retailer or learn more at [intel.com](https://www.intel.com).

This sample source code is released under the [Intel Sample Source Code License Agreement](#).

Intel and the Intel logo are trademarks of Intel Corporation in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.

Copyright © 2017, Intel Corporation. All rights reserved.

Review

- Do some review of concepts from the last lecture
- We will revisit kernel, stride, and pooling in the context of the Le-Net 5 model.


origin of CNN

LeNet-5

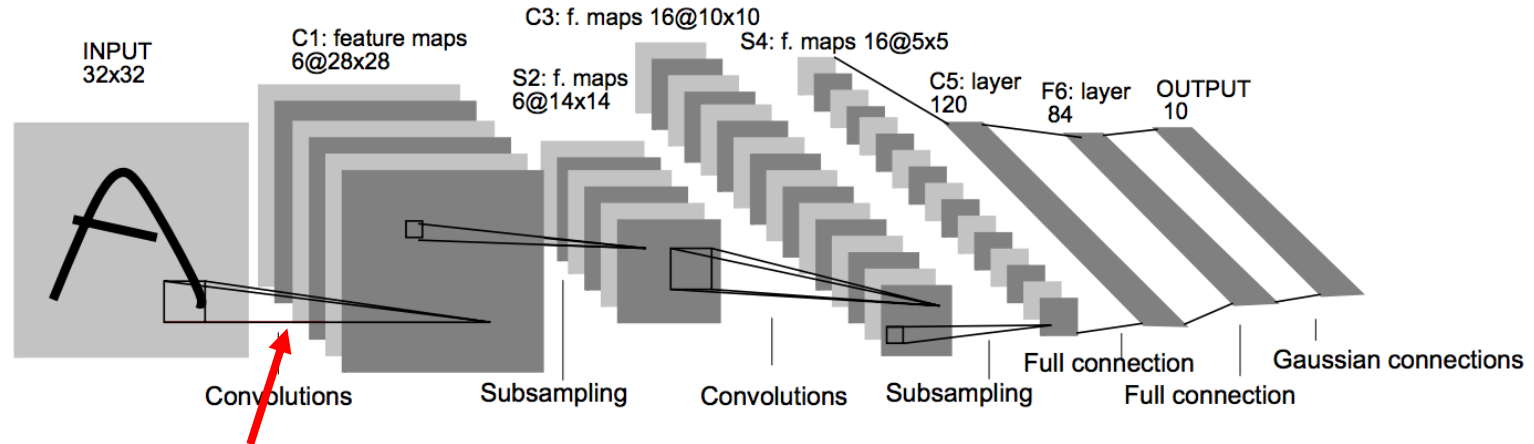
- Created by Yann LeCun in the 1990s
- Used on the MNIST data set.
- Novel Idea: Use convolutions to efficiently learn features on data set.

LeNet – Structure Diagram

Input: A 32 x 32 grayscale image
(28 x 28) with 2 pixels of padding
all around.

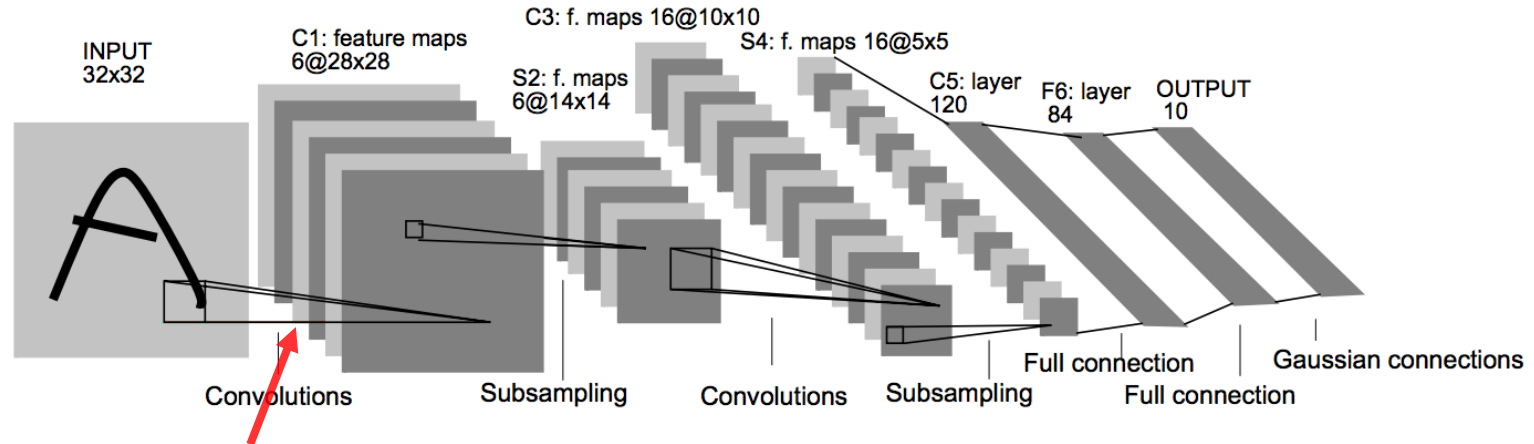


LeNet – Structure Diagram



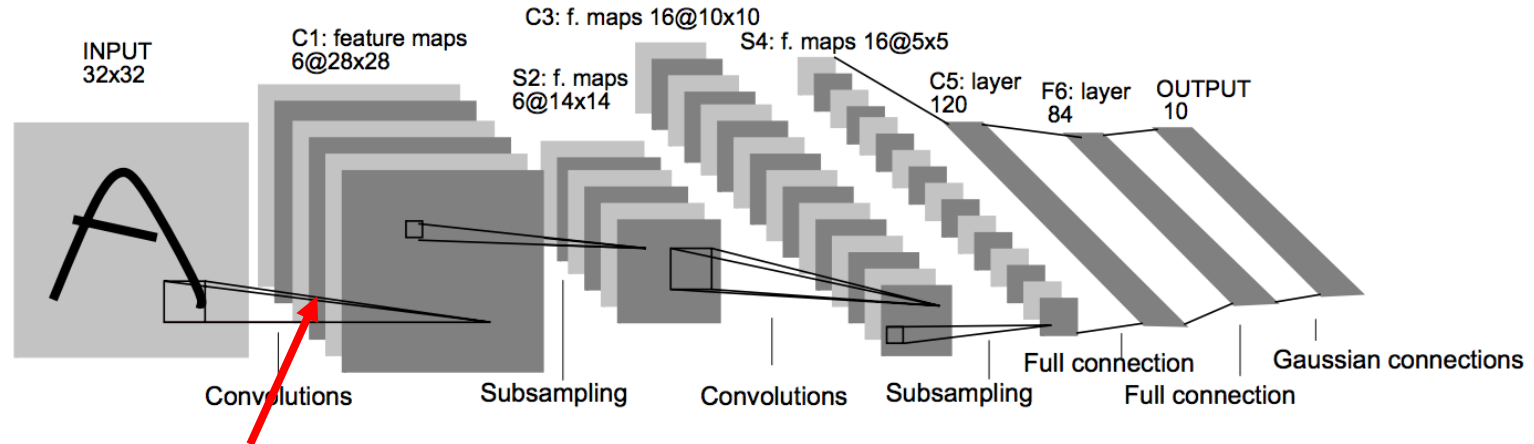
Next, we have a
convolutional layer.

LeNet – Structure Diagram



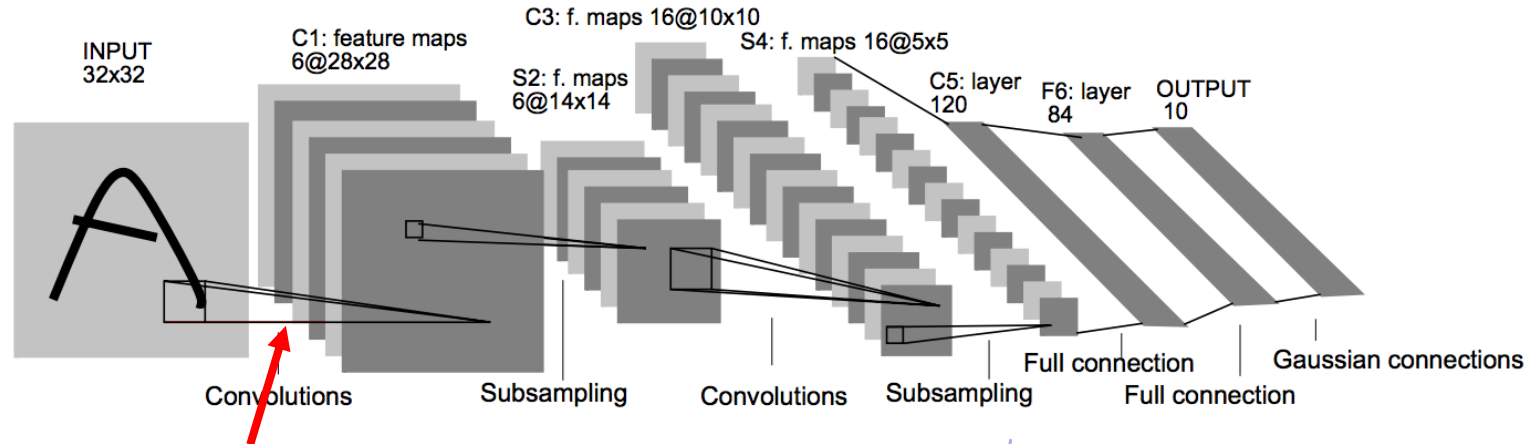
This is a 5x5 convolutional layer with stride 1.

LeNet – Structure Diagram



This means the resulting “filter” has dimension 28x28. (Why?)

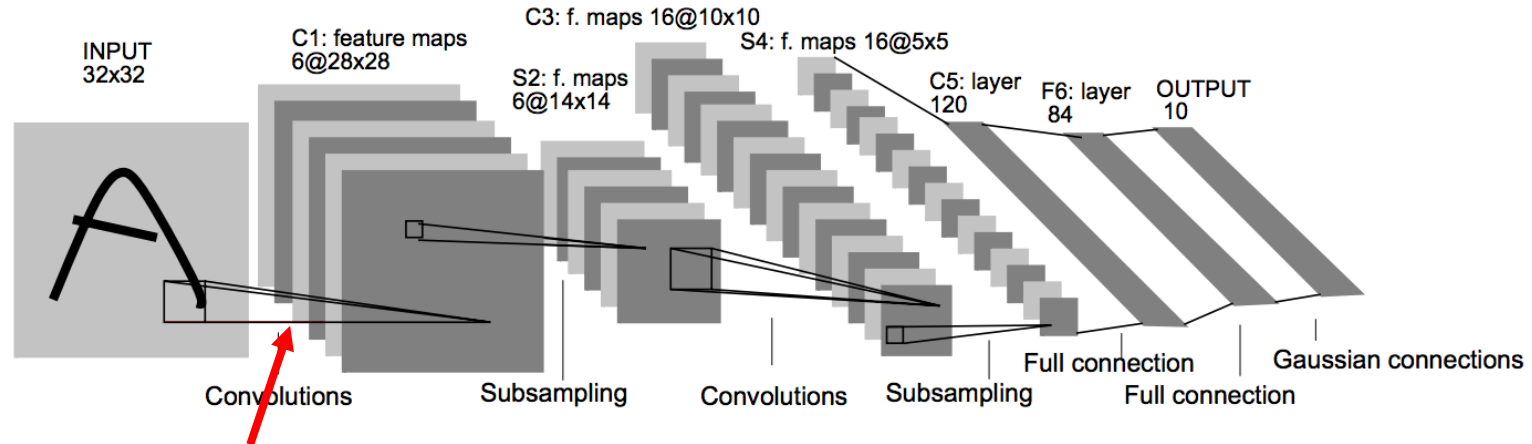
LeNet – Structure Diagram



They use a depth of 6. This means there are 6 different kernels that are learned.

high depth invariance design model
works on any behavior vs data


LeNet – Structure Diagram



They use a depth of 6. This means there are 6 different kernels that are learned.

So the output of this layer is 6x28x28.


LeNet – Structure Diagram



What is the total number of parameters in this layer?

$$\begin{aligned} ((6 \times 5) + 1) \times 6 &= 26 \times 6 \\ &= 156 \end{aligned}$$


LeNet – Structure Diagram



What is the total number of parameters in this layer?

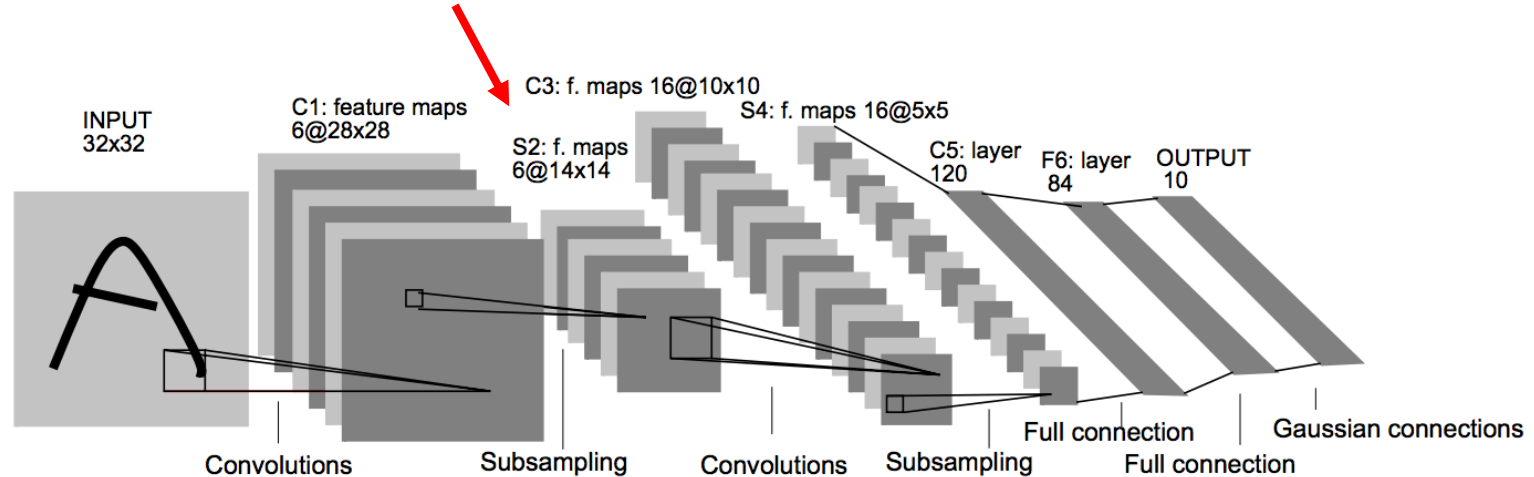
Answer: Each kernel has $5 \times 5 = 25$ weights (plus a bias term, so actually 26 parameters).
So total parameters = $6 \times 26 = 156$.

LeNet – Structure Diagram




LeNet – Structure Diagram

So output size is $6 \times 14 \times 14$ (we downsample by a factor of 2)



LeNet – Structure Diagram

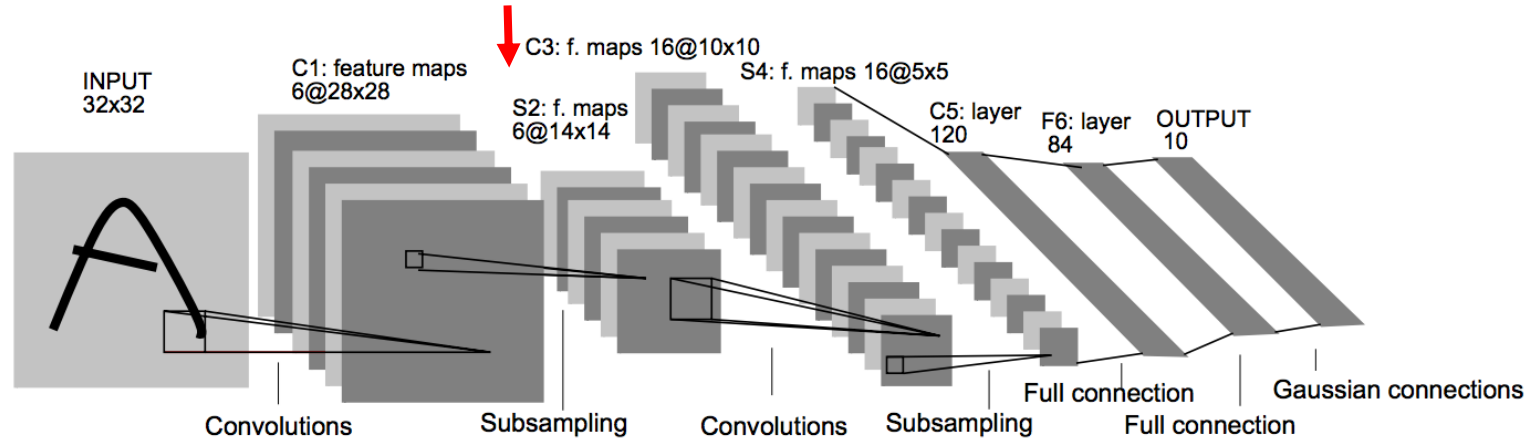
So output size is $6 \times 14 \times 14$
(we downsample by a factor of 2)



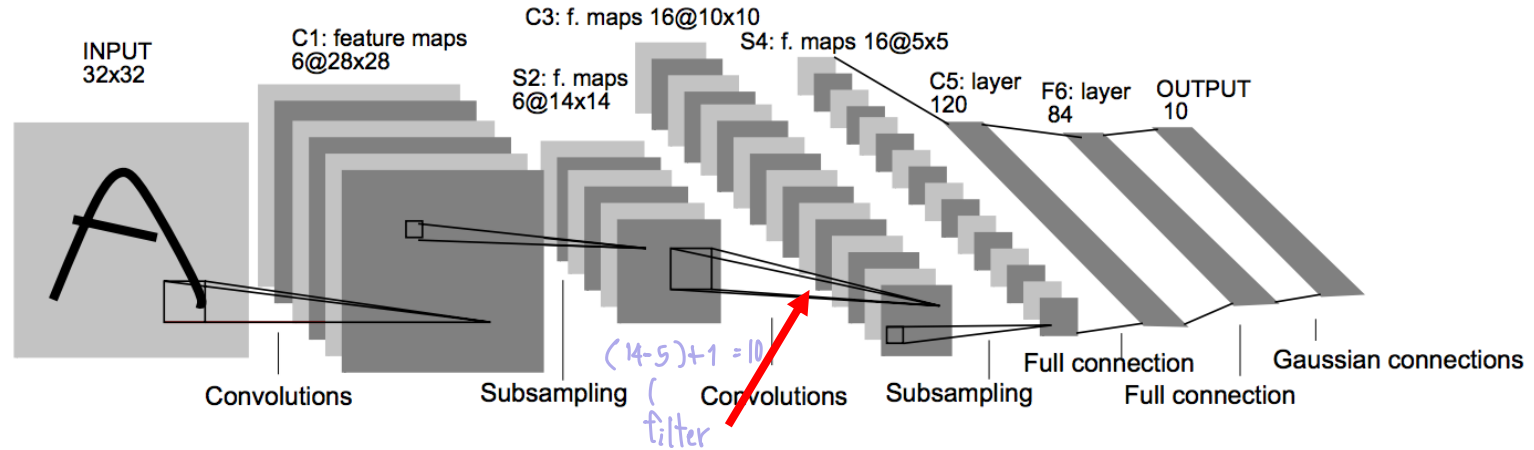
Note: The original paper actually does average pooling (rather than max pooling).

LeNet – Structure Diagram

No weights! (pooling layers have no weights to be learned – it is a fixed operation.)




LeNet – Structure Diagram



Another 5x5 convolutional layer with stride 1.
This time the depth is 16.

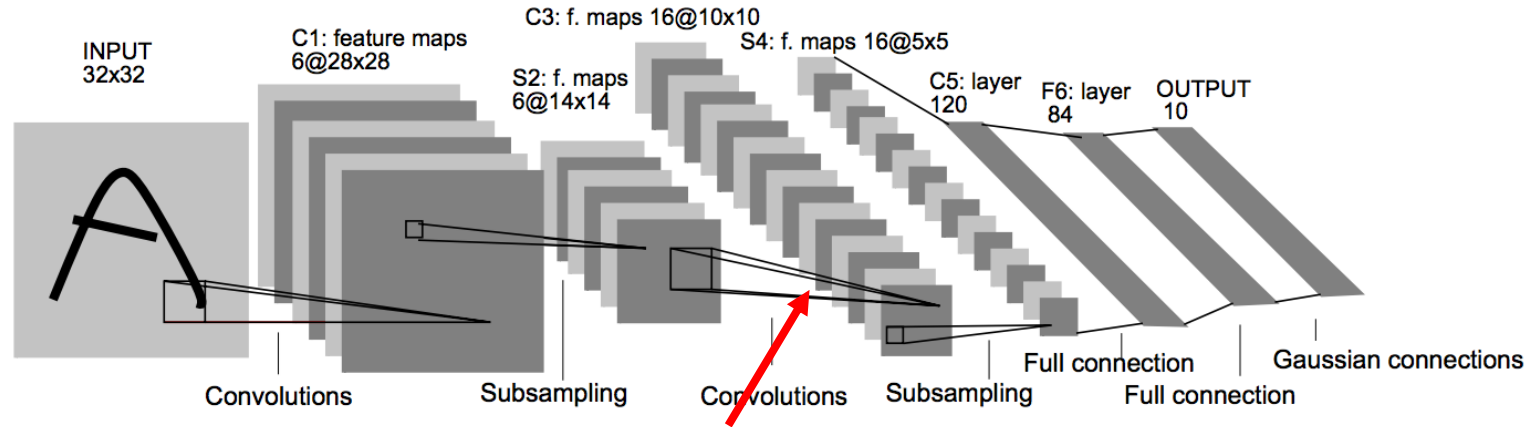
LeNet – Structure Diagram



Output size: 16 x 10 x 10
How many weights? (tricky!)

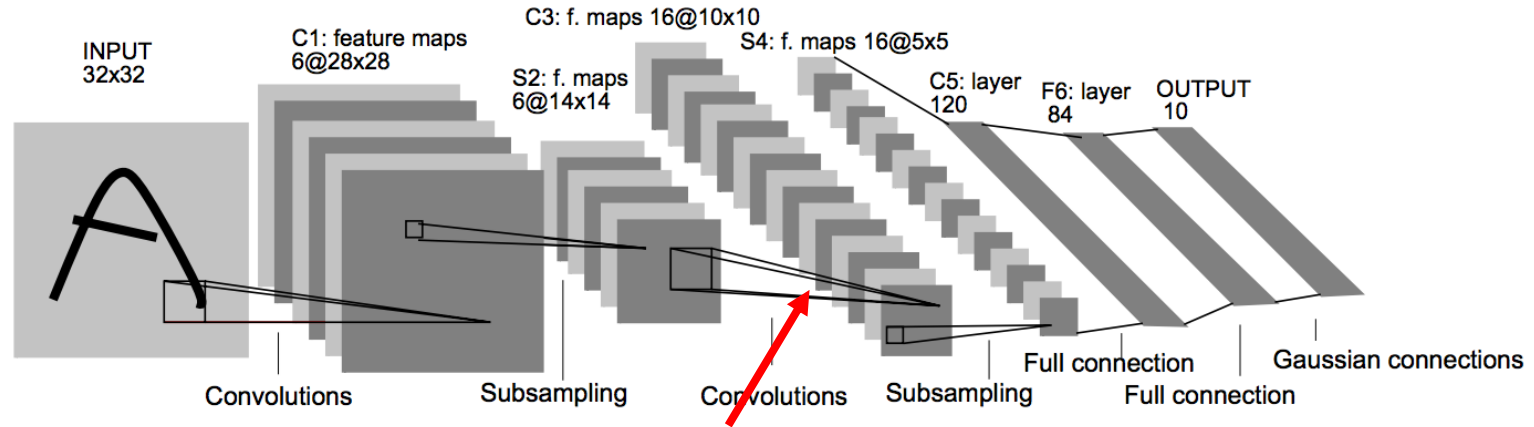
input
 $((6 \times 5 \times 5) + 1) \times 16$
 151×16

LeNet – Structure Diagram



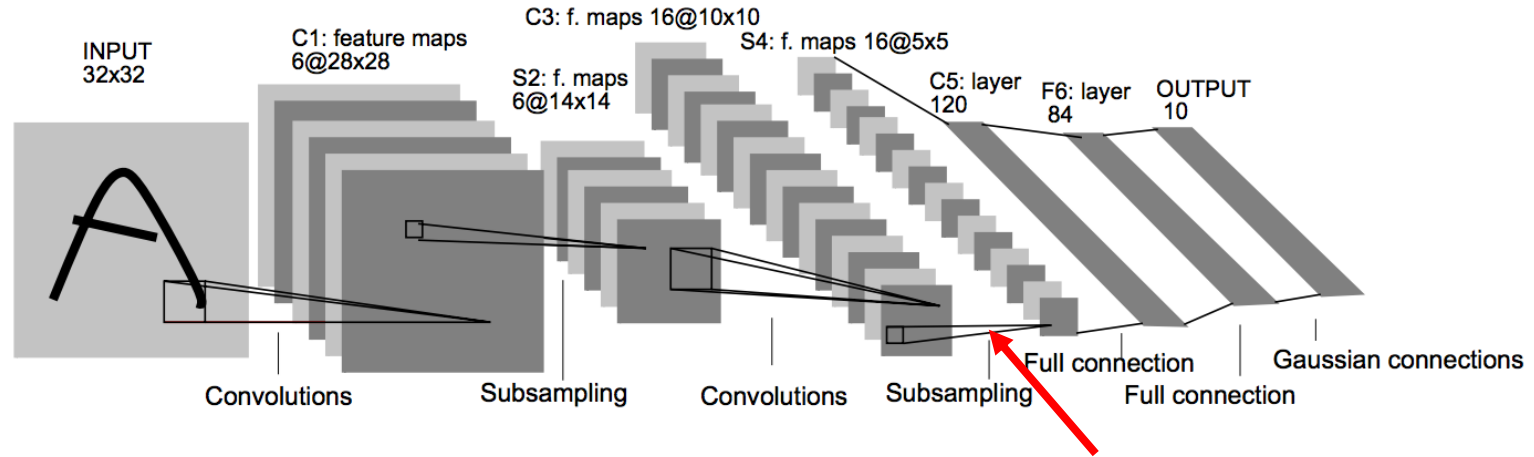
The kernels “take in” the full depth of the previous layer.
So each 5x5 kernel now “looks at” 6x5x5 pixels.
Each kernel has $6 \times 5 \times 5 = 150$ weights + bias term = 151.

LeNet – Structure Diagram



So, total weights for this layer = $16 * 151 = 2416$.


LeNet – Structure Diagram



Another 2x2 pooling layer.
Output is 16 x 5 x 5.

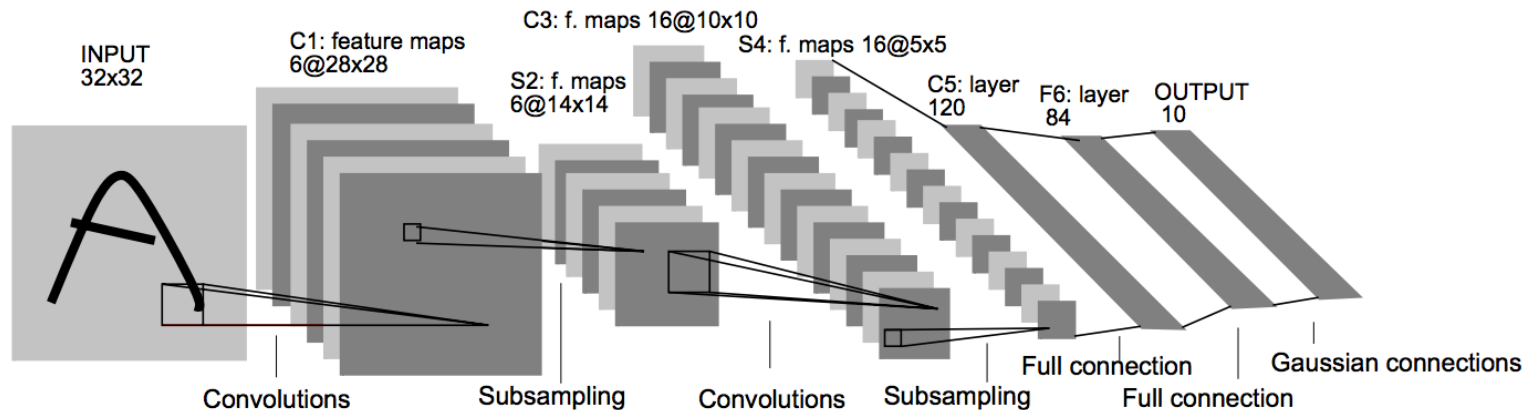
LeNet – Structure Diagram

We “flatten” this to a
length 400 vector (not shown)



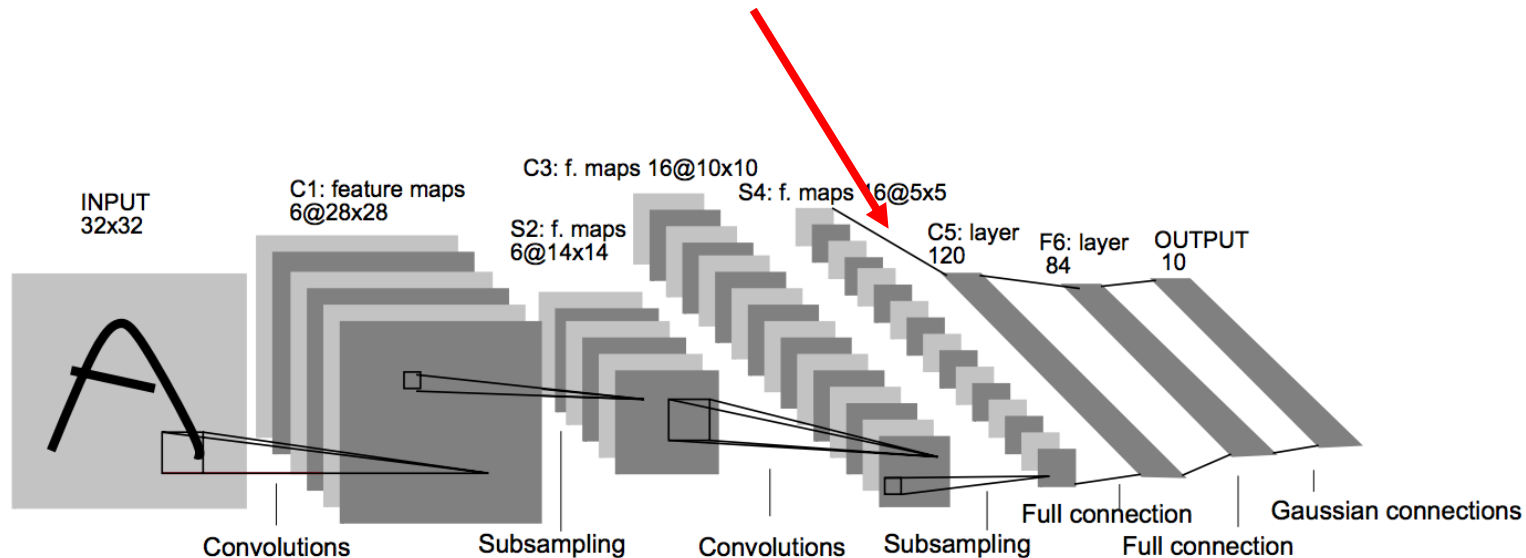
LeNet – Structure Diagram

The following layers are just fully connected layers!

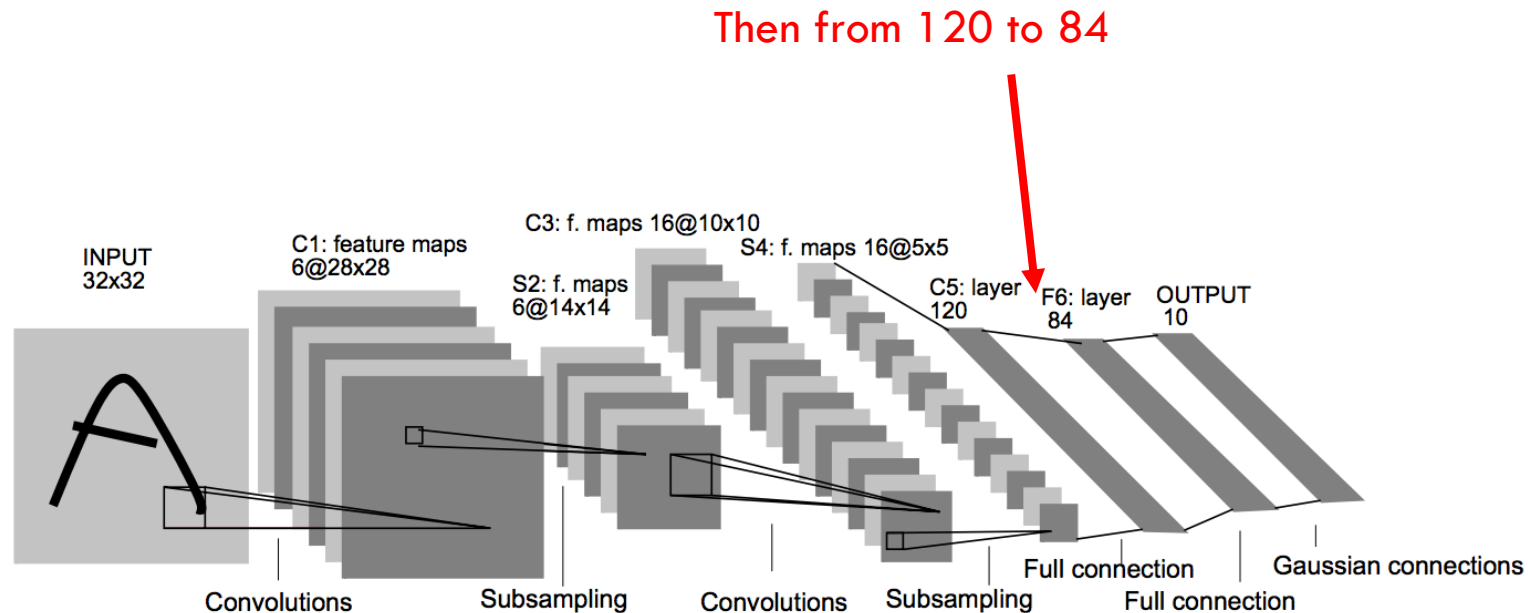


LeNet – Structure Diagram

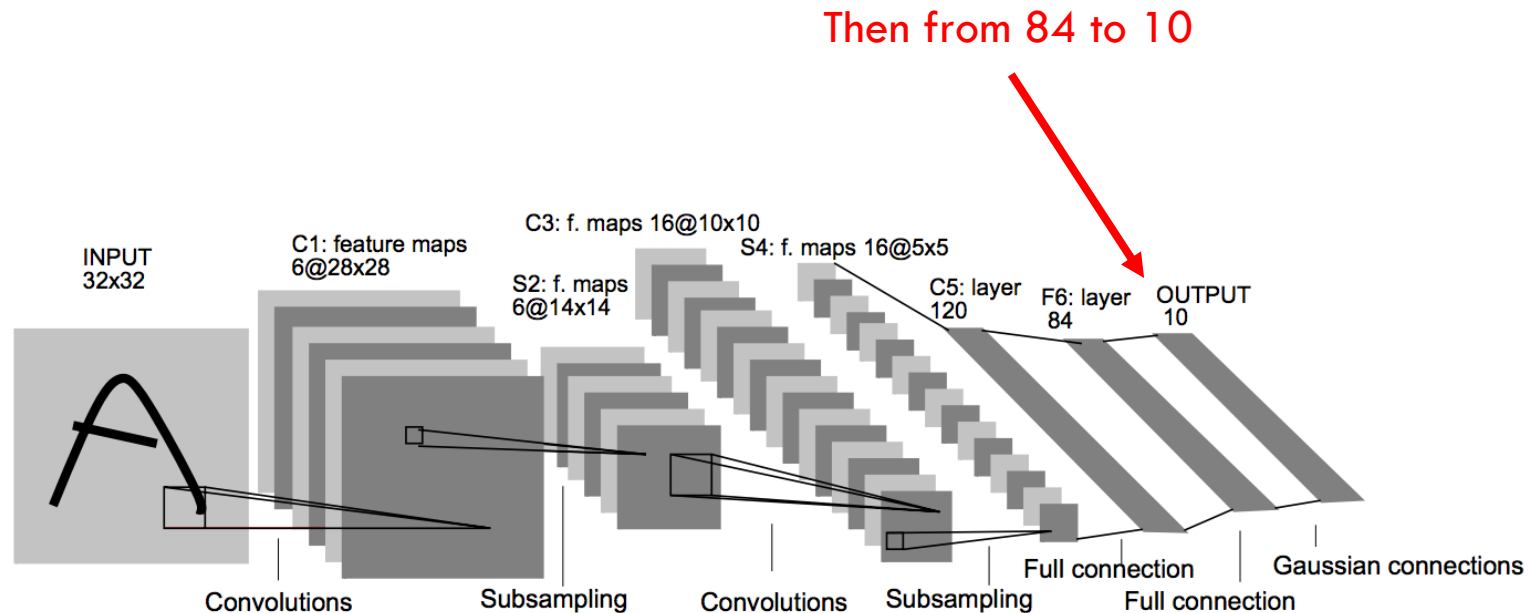
From 400 to 120



LeNet – Structure Diagram




LeNet – Structure Diagram



LeNet – Structure Diagram

And a softmax output
of size 10 for the 10
digits



LeNet-5

How many total weights in the network?

$$\text{Conv1: } 6 * 1 * 5 * 5 + 6 = 156$$

$$\text{Conv3: } 16 * 6 * 5 * 5 + 16 = 2,416$$

$$\text{FC1: } 400 * 120 + 120 = 48,120$$

$$\text{FC2: } 120 * 84 + 84 = 10,164$$

$$\text{FC3: } 84 * 10 + 10 = 850$$

$$\text{Total: } \text{(output 10 class)} = 61,706$$


Less than a single FC layer with [1 200x1 200] weights!

Note that Convolutional Layers have relatively few weights.

Transfer Learning

Motivation

- Early layers in a Neural Network are the hardest (i.e. slowest) to train
- Due to **vanishing gradient** property
- But these "primitive" features should be general across many image classification tasks
- A **feature map** is the output of a **convolutional layer** representing **specific features** in the input image or feature map.



Motivation

- Later layers in the network are capturing features that are more particular to the specific image classification problem.
- Later layers are easier (quicker) to train since adjusting their weights has a more immediate impact on the final result. *from gradient formula, back propagating's sigmoid make the earlier weight smaller and smaller*


Motivation

- Famous, Competition-Winning Models are difficult to train from scratch
 - Huge datasets (like ImageNet)
 - Long number of training iterations
 - Very heavy computing machinery
 - Time experimenting to get hyper-parameters just right

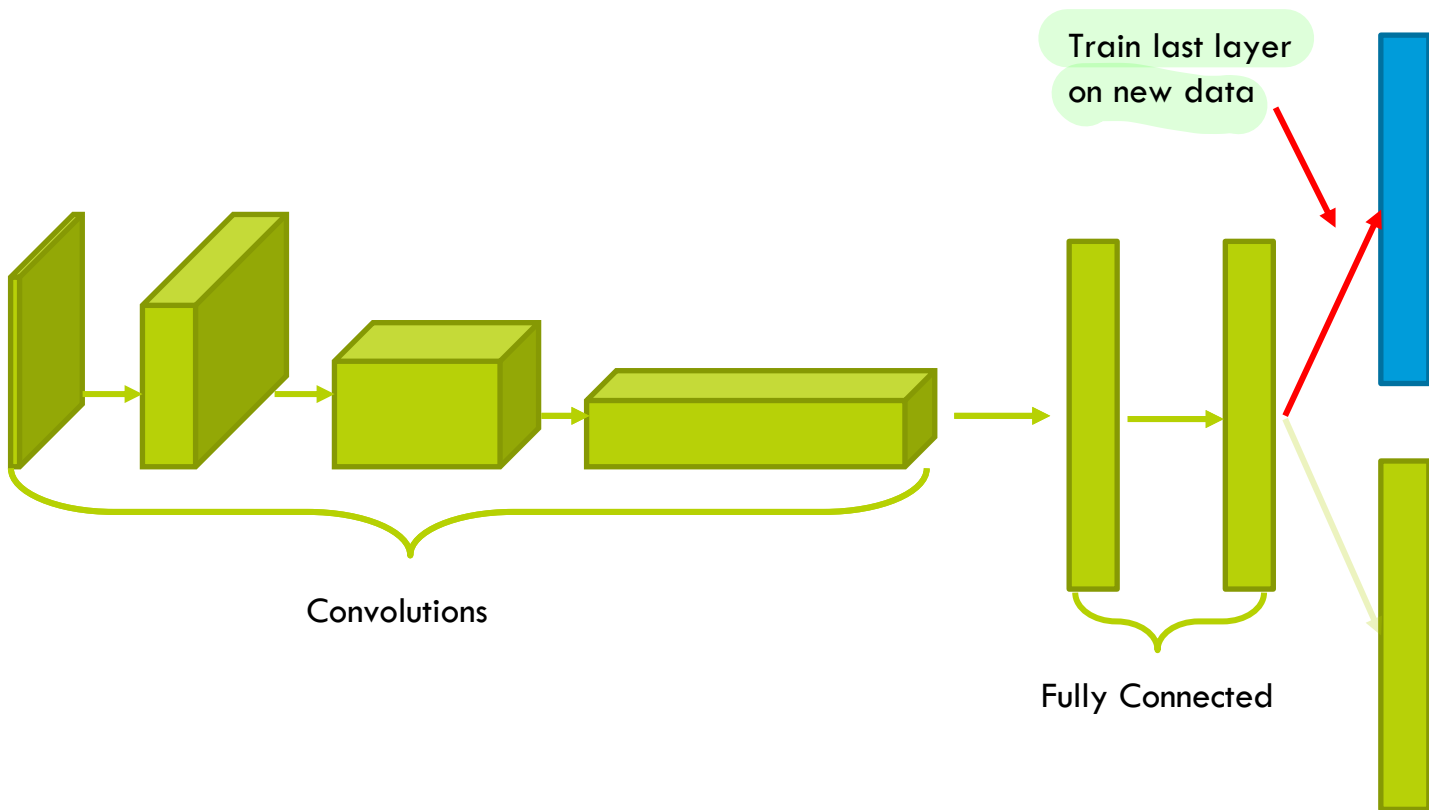
Transfer Learning

- However, the basic features (edges, shapes) learned in the early layers of the network *should* generalize.
- Results of the training are just weights (numbers) that are easy to store.
- Idea: keep the early layers of a pre-trained network, and re-train the later layers for a specific application
- This is called **Transfer Learning**.

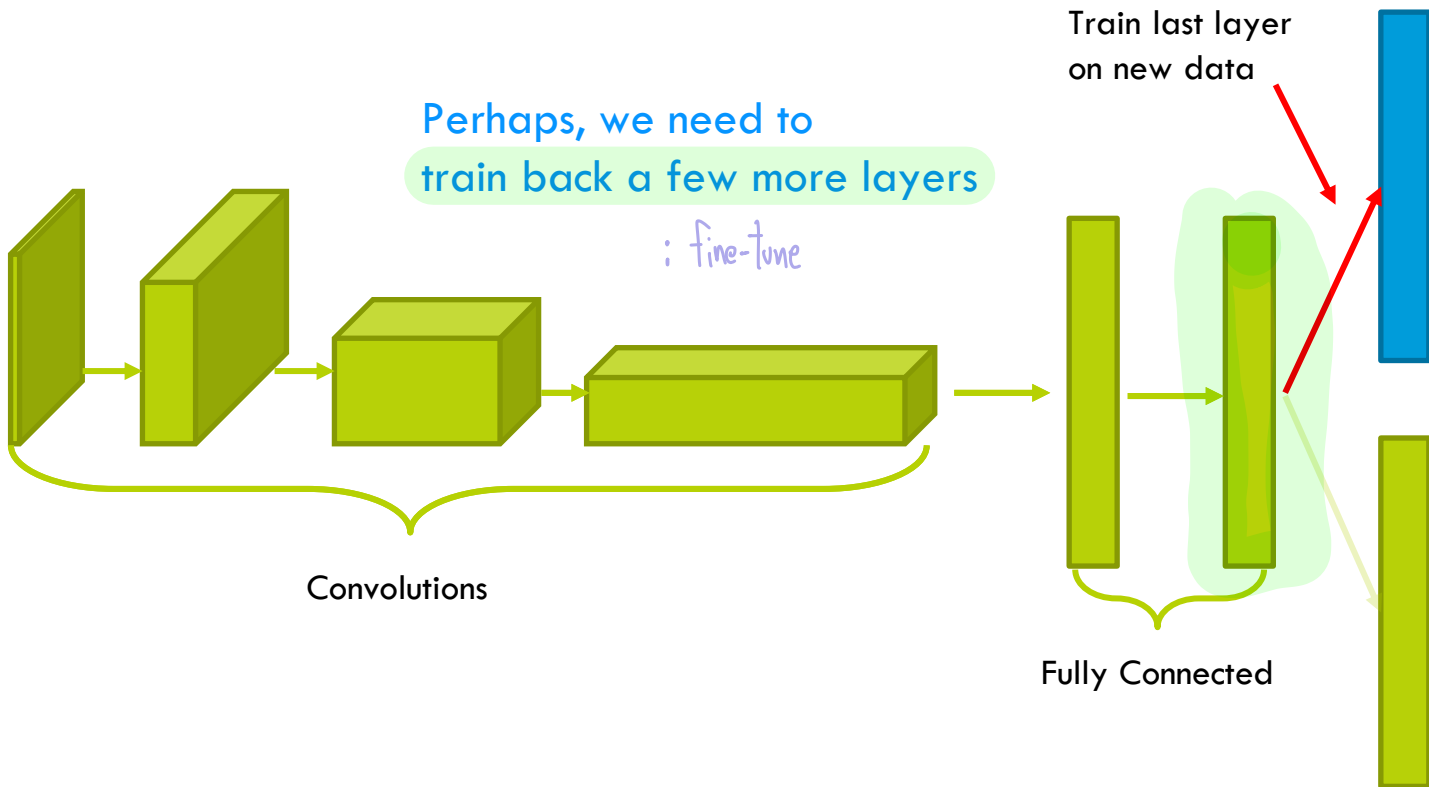
Transfer Learning



Transfer Learning



Transfer Learning



Transfer Learning Options

- The additional training of a pre-trained network on a specific new dataset is referred to as “Fine-Tuning”
- There are different options on “how much” and “how far back” to fine-tune.
 - Should I train just the very last layer?
 - Go back a few layers?
 - Re-train the entire network (from the starting point of the existing network)?

Guiding Principles for Fine-Tuning




- While there are no “hard and fast” rules, there are some guiding principles to keep in mind.
 - 1) The more similar your data and problem are to the source data of the pre-trained network, the less fine-tuning is necessary.
 - E.g. Using a network trained on ImageNet to distinguish “dogs” from “cats” should need relatively little fine-tuning. It already distinguished different breeds of dogs and cats, so likely has all the features you will need.

Guiding Principles for Fine-Tuning

- 2) The more data you have about your specific problem, the more the network will benefit from longer and deeper fine-tuning.
- E.g. If you have only 100 dogs and 100 cats in your training data, you probably want to do very little fine-tuning. If you have 10,000 dogs and 10,000 cats you may get more value from longer and deeper fine-tuning.

Guiding Principles for Fine-Tuning

- 3) If your data is substantially different in nature than the data the source model was trained on, Transfer Learning may be of little value.
- E.g. A network that was trained on recognizing typed Latin alphabet characters would not be useful in distinguishing cats from dogs. But it likely would be useful as a starting point for recognizing Cyrillic alphabet characters.



Software