Software

Review of Machine Learning

Materials from

- Intel Deep Learning https://www.intel.com/content/www/us/en/developer/learn/course-deep-learning.html
- Introduction to Neural Networks https://www.deeplearning.ai/

Legal Notices and Disclaimers

Why is Deep Learning Taking Off?

Deep learning is taking off due to a large amount of data available through the digitization of the society, faster computation and innovation in the development of neural network algorithm.

Scale drives deep learning progress

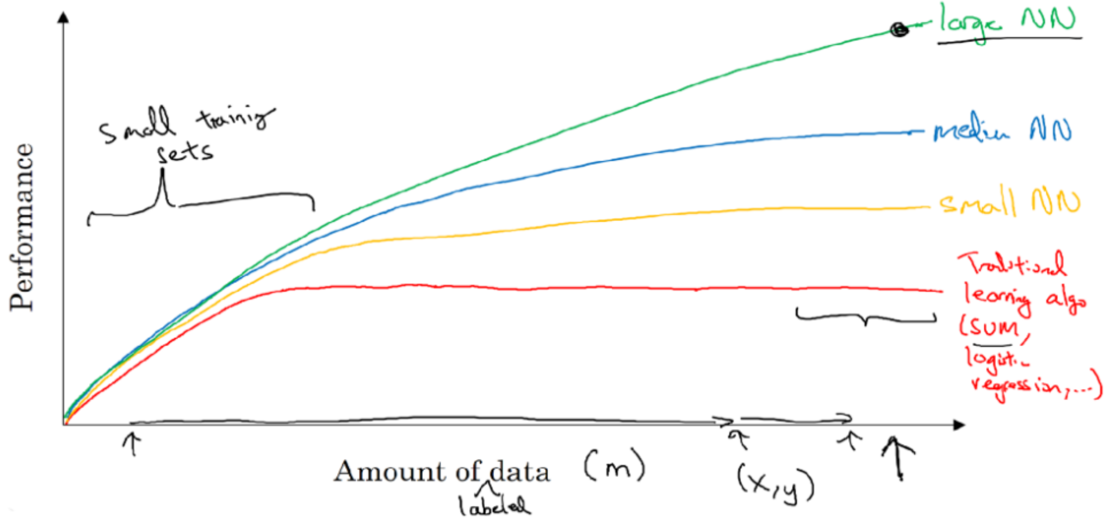

Two things have to be considered to get to the high level of performance:

1. Being able to train a big enough neural network
2. Huge amount of labeled data

Process of Training a Neural Network

The process of training a neural network is iterative.

- Focus on non linear equation but is inspired by the basic linear one
-

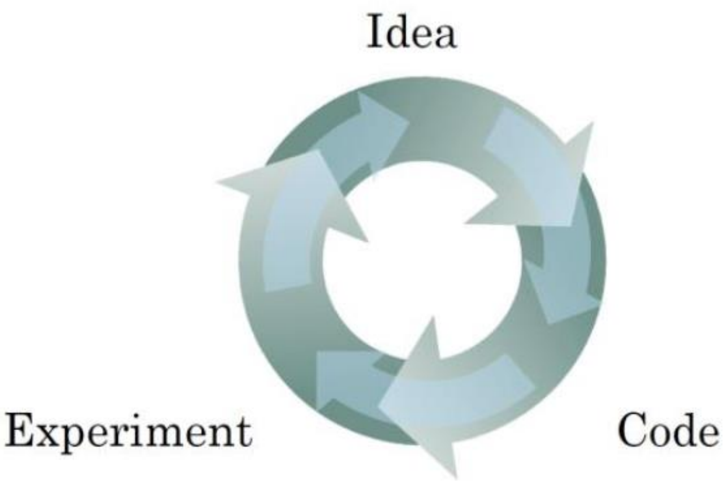

It could take a good amount of time to train a neural network, which affects your productivity. Faster computation helps to iterate and improve new algorithm.

What is Machine Learning?

Machine learning allows

The model and equation

computers to learn and infer from data.

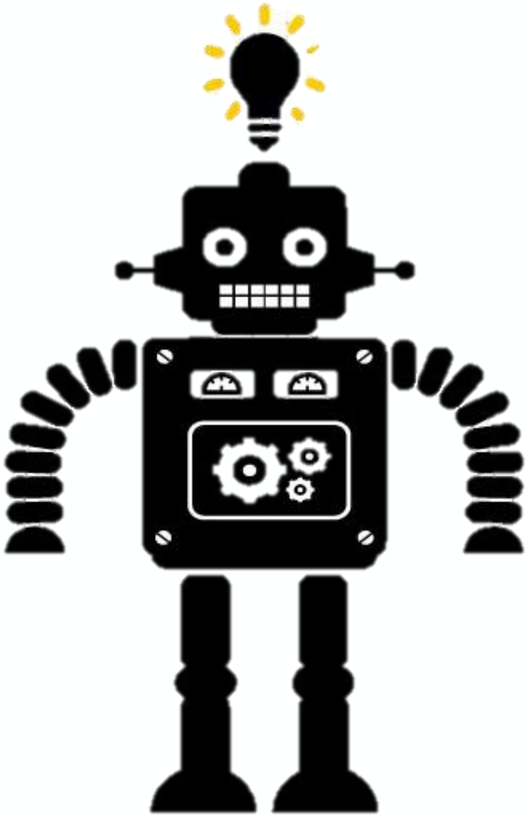

Classical Programming and Machine Learning

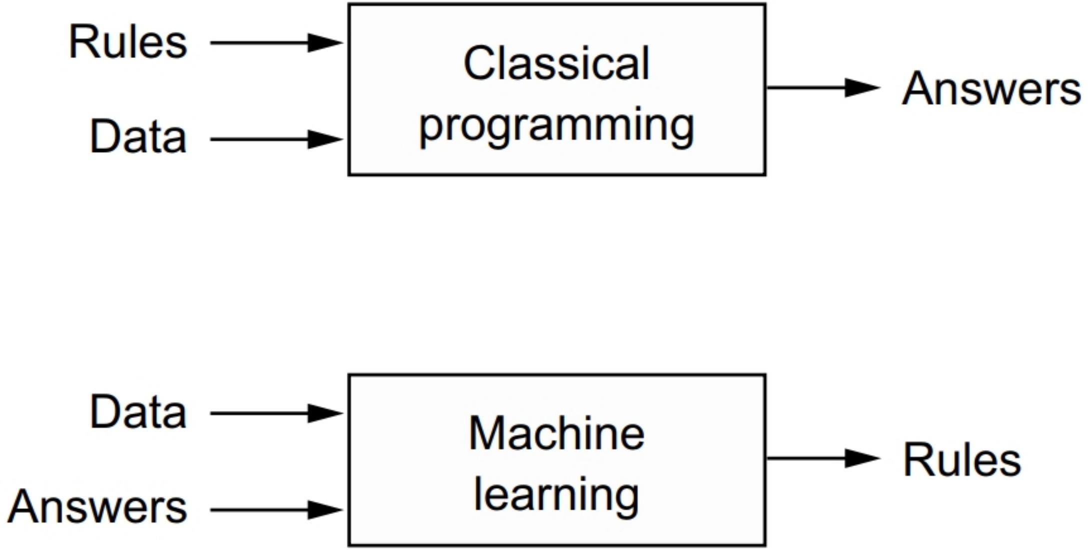

Image Source:
Deep Learning with Python, Second Edition
By Francois Chollet

Artificial Intelligence, Machine Learning, and Deep Learning

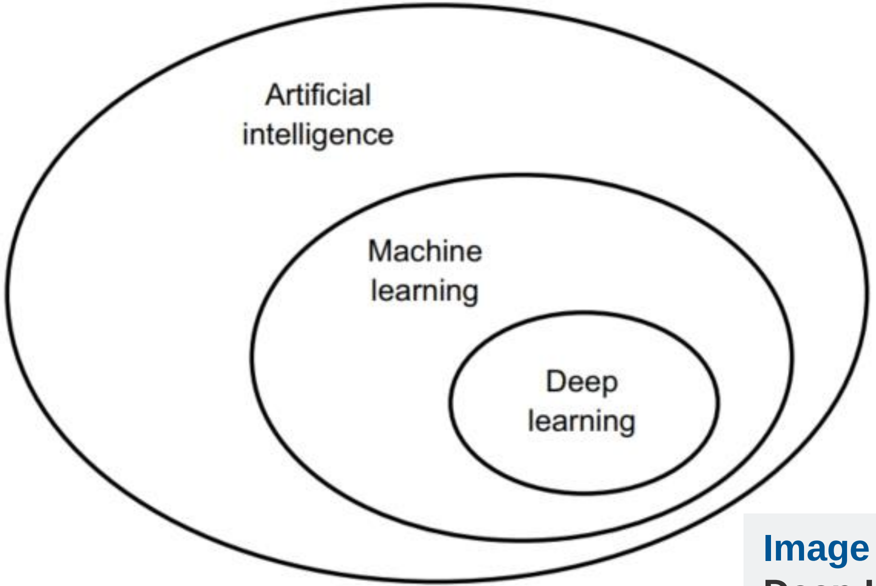

Image Source:
Deep Learning with Python, Second Edition
By Francois Chollet

Types of Machine Learning

Supervised — data points have known outcome

Data with label (y)

Unsupervised — data points have unknown outcome

for Grouping purpose

Types of Supervised Learning

the outcome result in mumerical format, like the colories calculated from …

Regression outcome is continuous (numerical)

Classification outcome is a category

Machine Learning Vocabulary

- **Target:** predicted category or value of the data (column to predict)
- **Features:** properties of the data used for prediction (non-target columns)
- **Example:** a single data point within the data (one row) record
- **Label:** the target value for a single data point

Machine Learning Vocabulary (Synonyms)

- **Target:** Response, Output, Dependent Variable, Labels
- **Features:** Predictors, Input, Independent Variables, Attributes
- **Example:** Observation, Record, Instance, Datapoint, Row
- **Label:** Answer, y-value, Category

Supervised Learning Overview

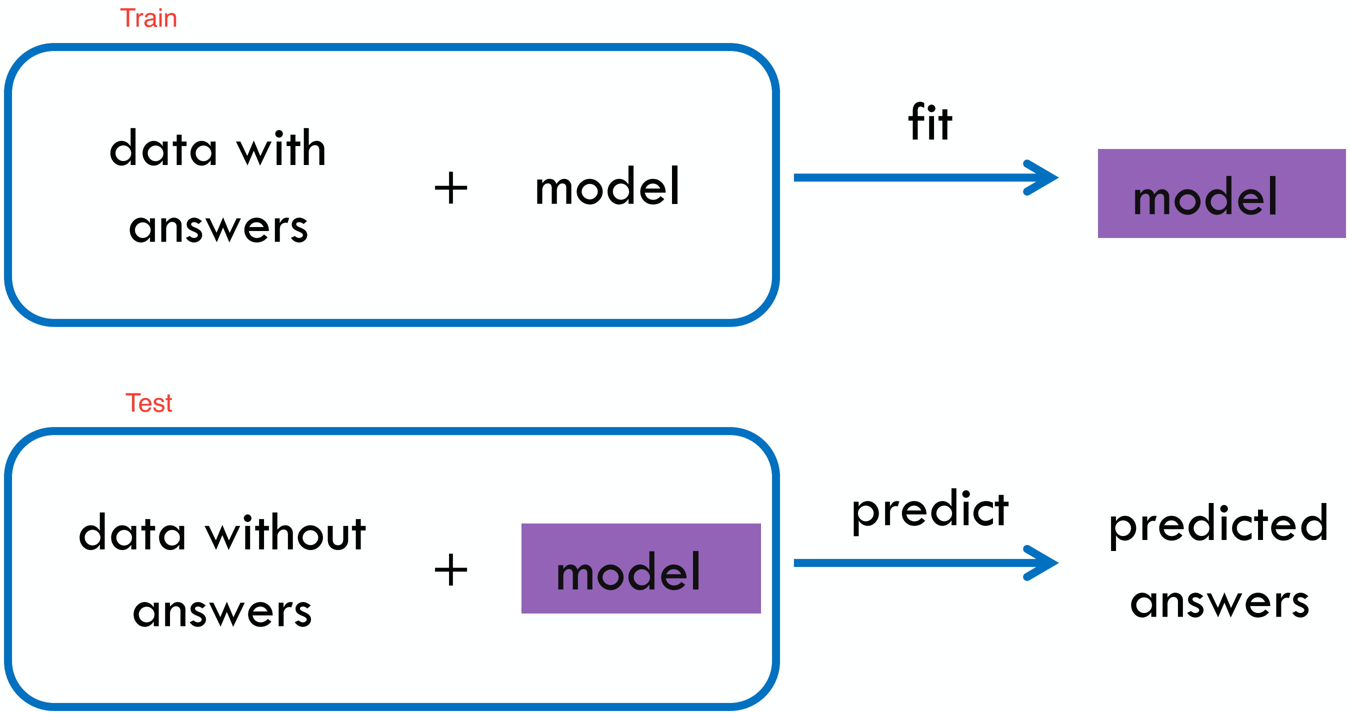

Regression: Numeric Answers

movie data with revenue + model $\xrightarrow{\text{fit}}$ model

movie data (unknown revenue) + model $\xrightarrow{\text{predict}}$ predicted revenue

Classification: Categorical Answers

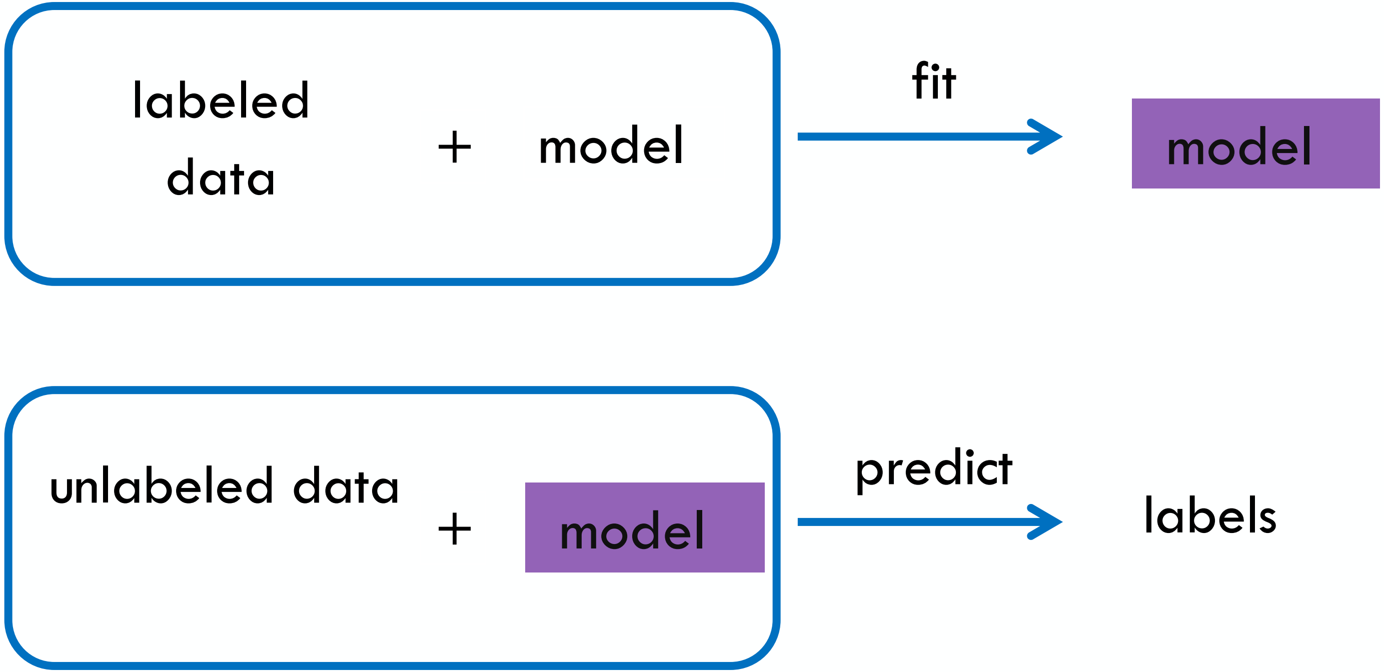

Classification: Categorical Answers

emails labeled as spam/not spam + model $\xrightarrow{\text{fit}}$ model

unlabeled emails + model $\xrightarrow{\text{predict}}$ spam or not spam

Three Types of Classification Predictions

- **Hard Prediction:** Predict a single category for each instance.
- **Ranking Prediction:** Rank the instances from most likely to least likely. (binary classification)
- **Probability Prediction:** Assign a probability distribution across the classes to each instance.

Metrics for Classification

- **Hard Prediction:** Accuracy, Precision, Recall (Sensitivity), Specificity, F1 Score
- **Ranking Prediction:** AUC (ROC), Precision-Recall Curves

  Area Under Curve
- **Probability Prediction:** Log-loss (aka Cross-Entropy), Brier Score

Metrics for Regression

- **Root Mean Square Error (RMSE)** ignore the negative and positive term so we use “square”

$$RMSE = \sqrt{\frac{1}{n}\sum_{i=1}^{n}(y_i - \hat{y}_i)^2}$$

Y hat refer to the predicted value
so if the prediction is correct,
“ y - (y hat) ” should equal to 0

- **Mean Absolute Error (MAE)**

$$MAE = \frac{1}{n}\sum_{i=1}^{n}|y_i - \hat{y}_i|$$

Fitting Training and Test Data

To prevent overfitting, we expect the model to be able to predict the new upcoming data so we split the whole data to the training one and the test one

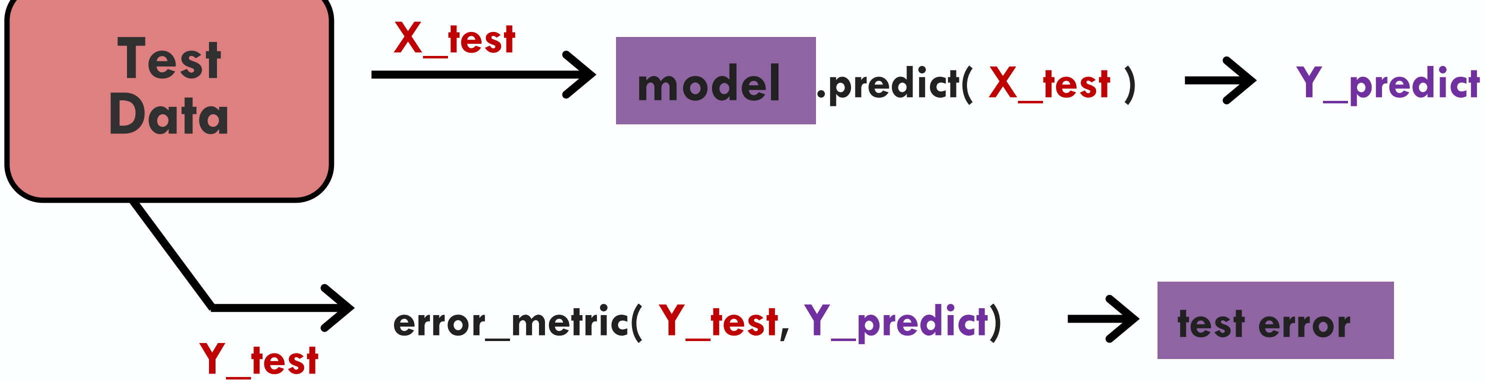

Using Training and Test Data

Training Data

fit the model

Test Data

measure performance

- predict label with model
- compare with actual value
- measure error

Using Training and Test Data

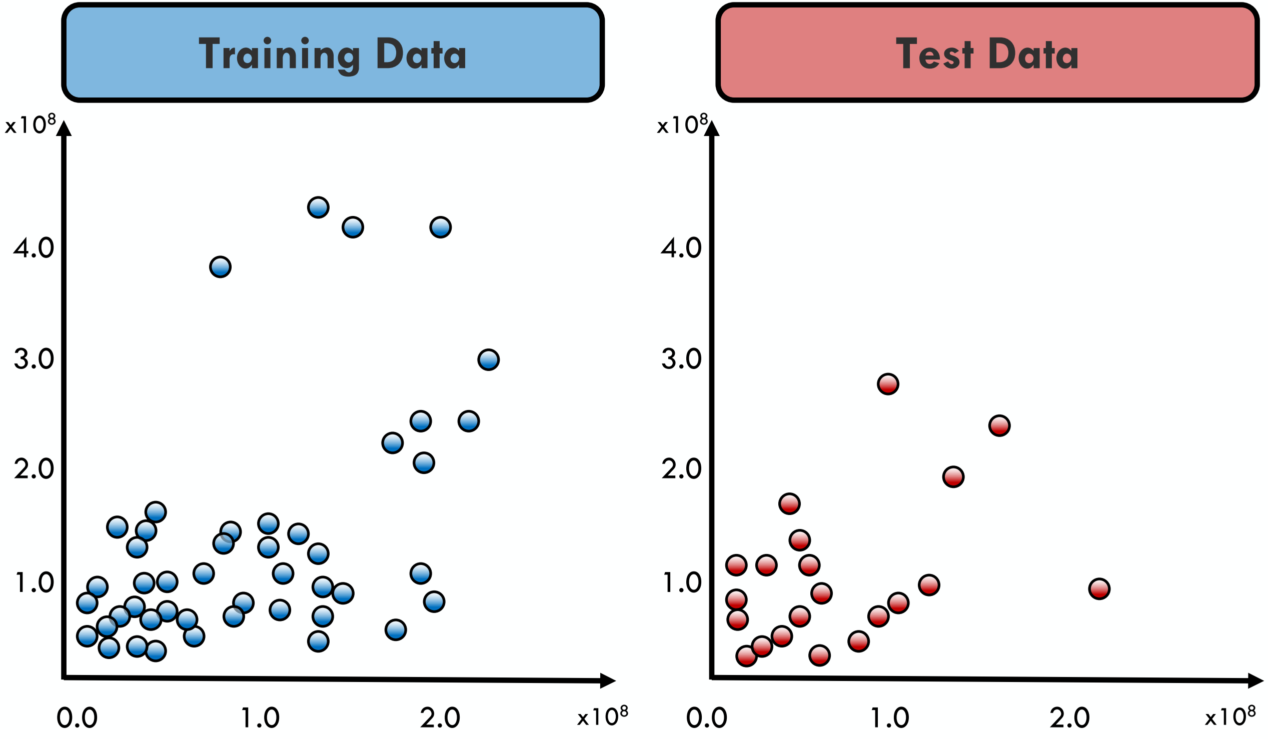

Using Training and Test Data

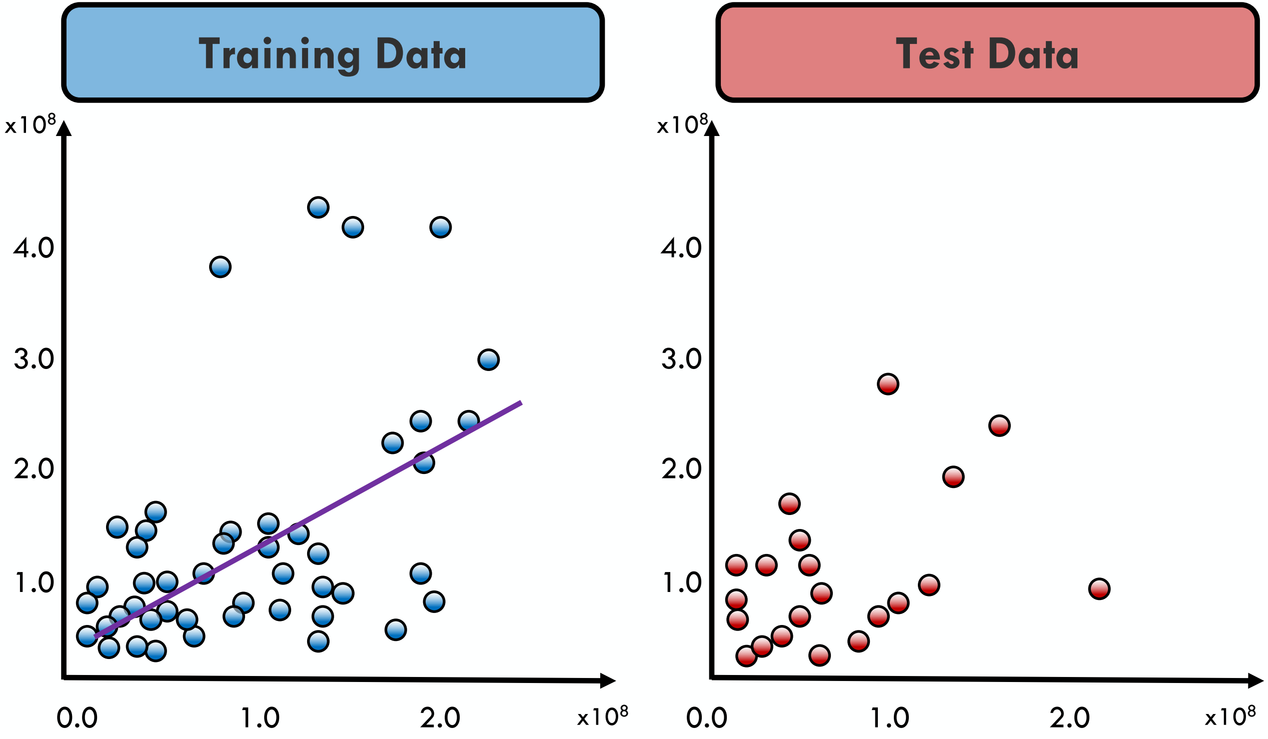

Fit the model

Using Training and Test Data

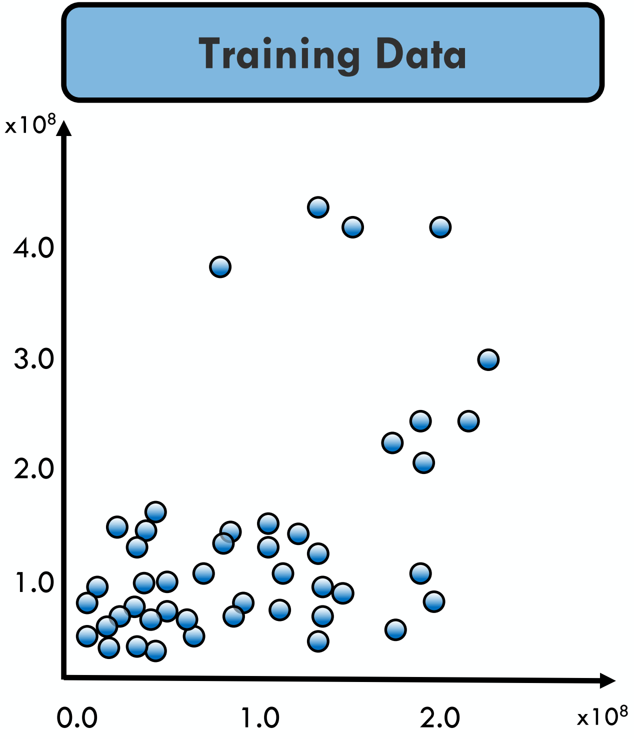

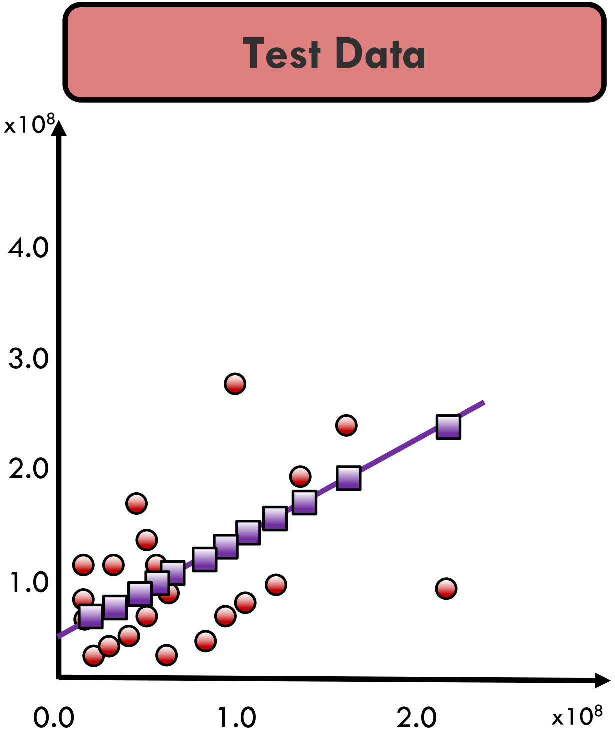

Make predictions

Using Training and Test Data

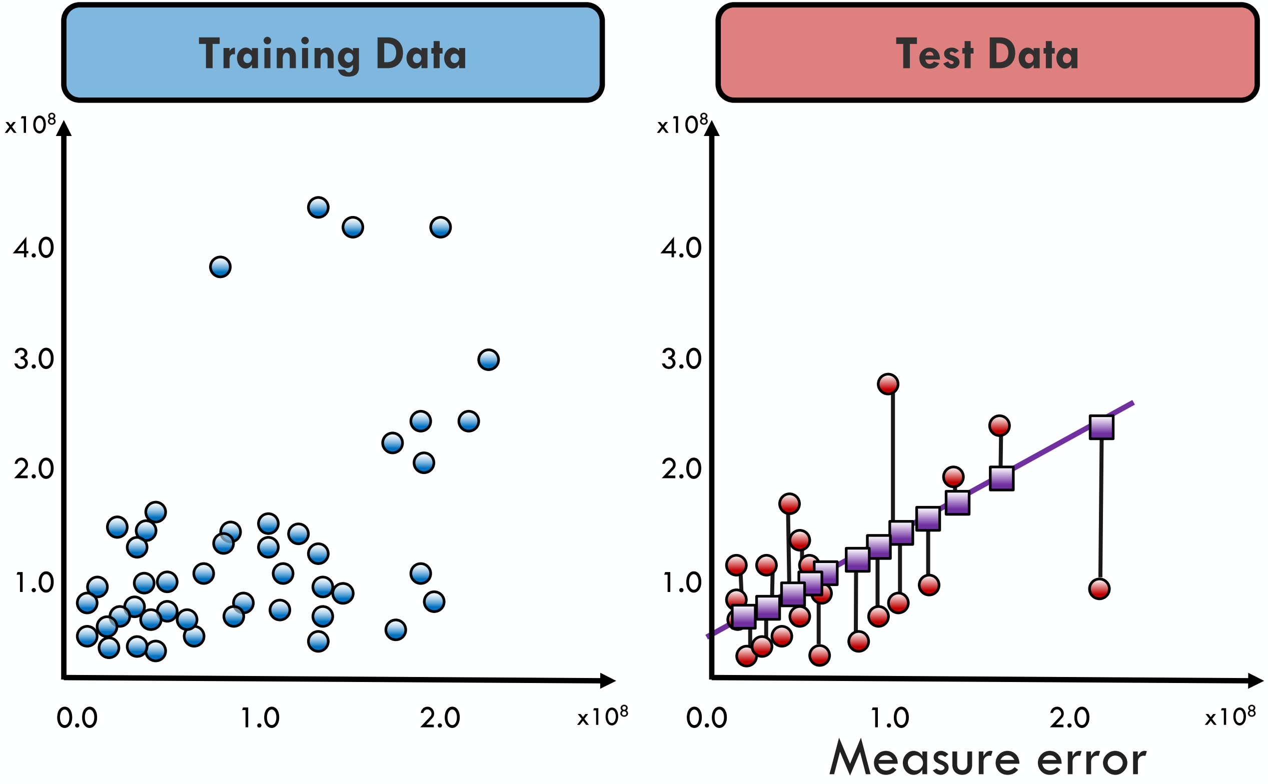

How Well Does the Model Generalize?

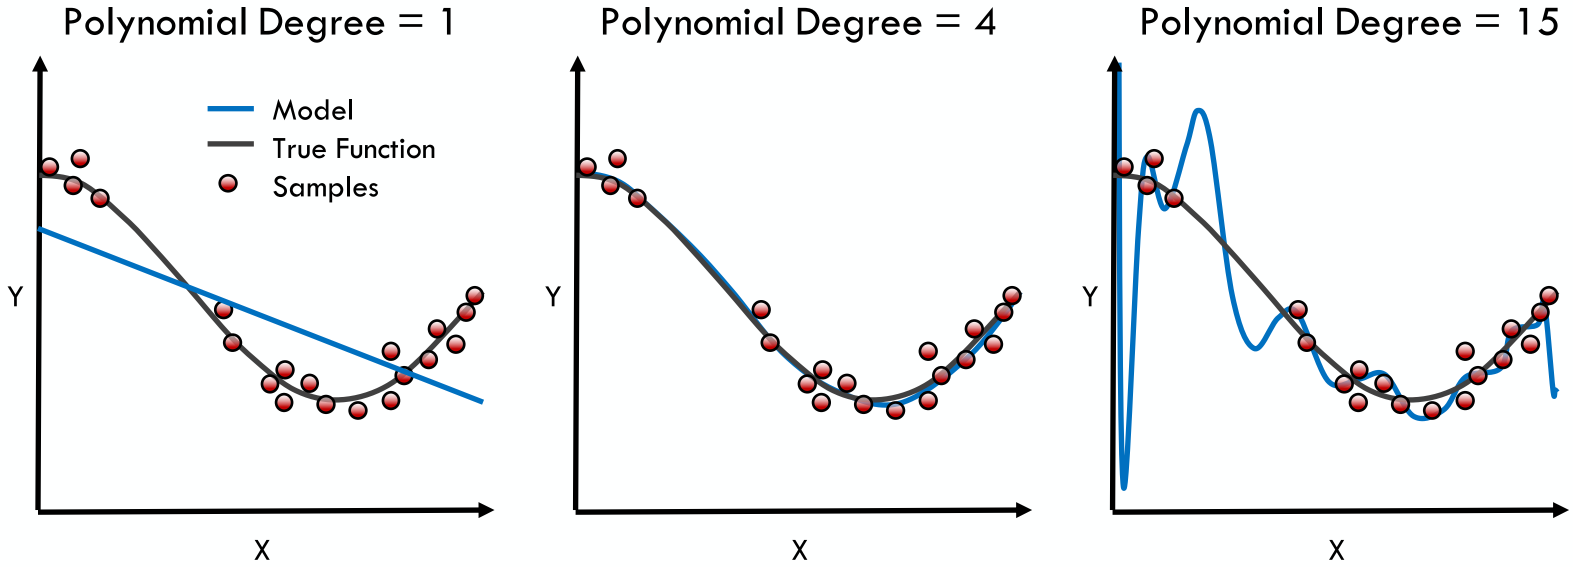

Poor on Training Set
Poor at Predicting

Just Right

Very Good on Training Set
Poor at Predicting

Underfitting vs Overfitting

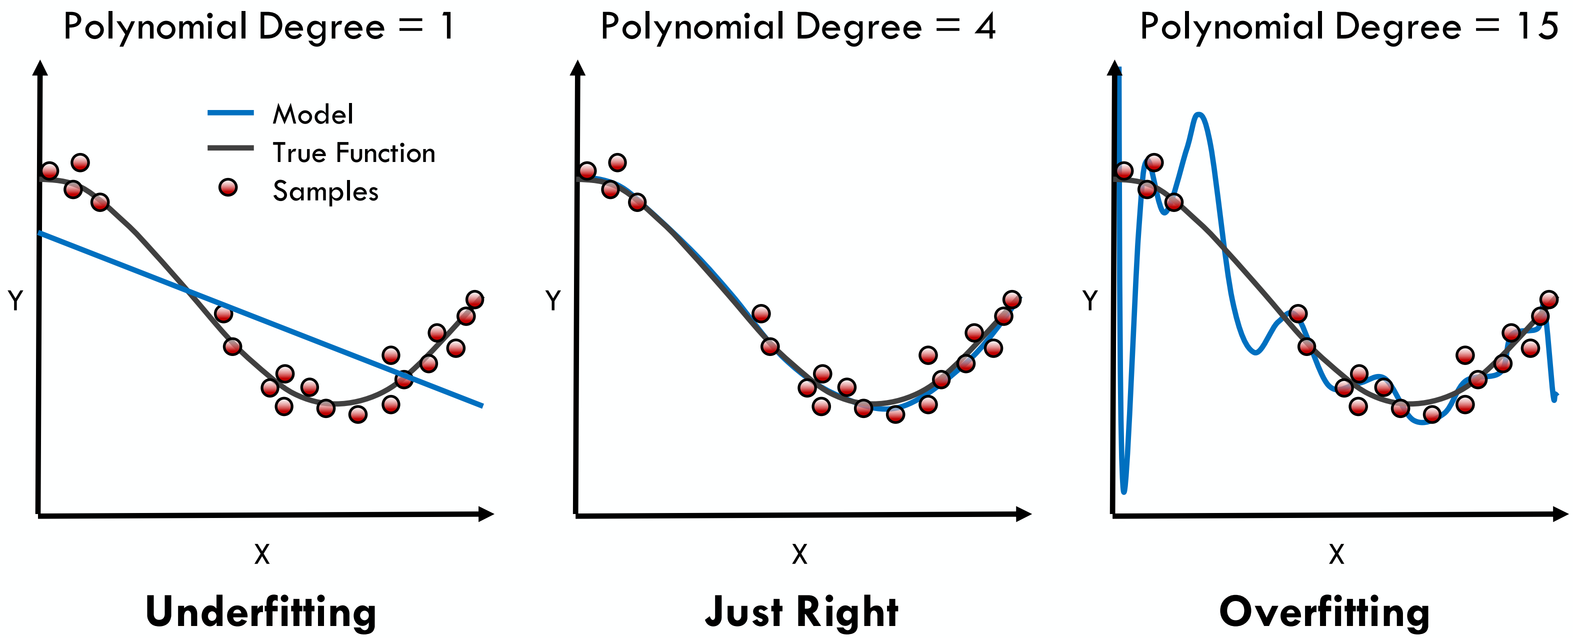

Bias – Variance Tradeoff

Bias : inablitity of the model to learn from the data
Variance: the different in fits between data sets

Low bias : learning well with the training data
High variance refers to the high difference of fitting between the data sets (more than 2 sets)

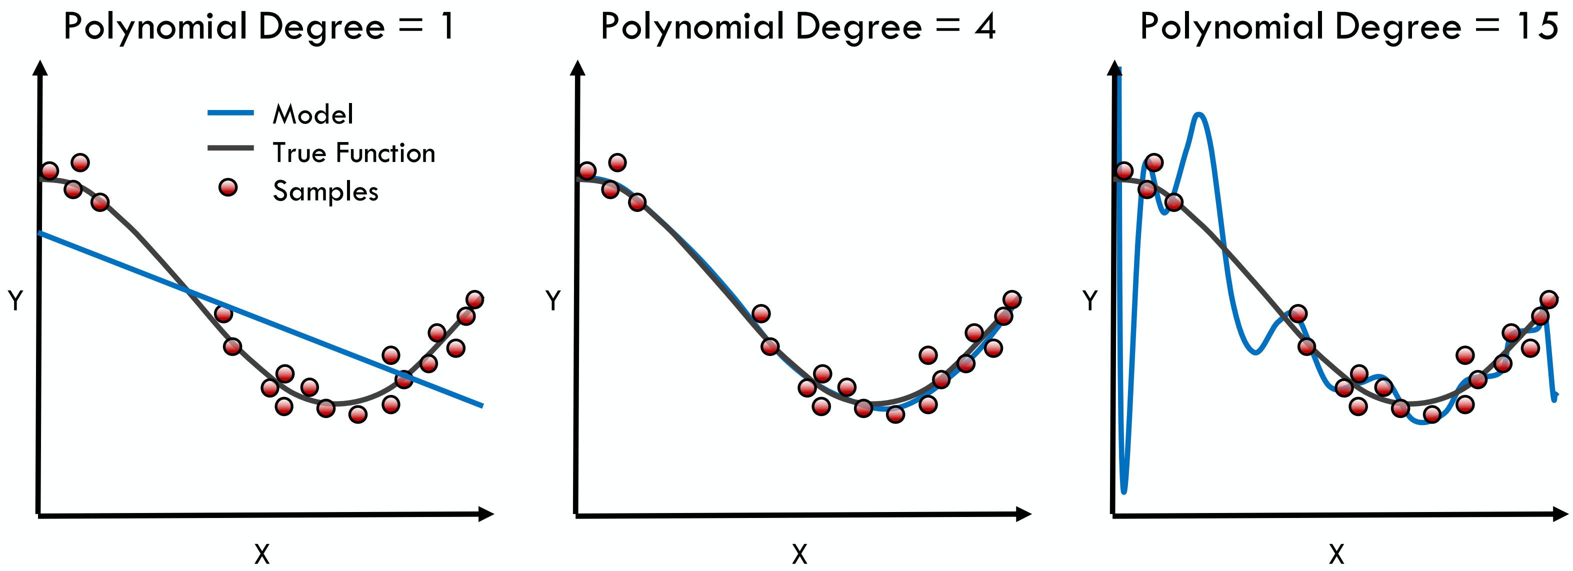

High Bias
Low Variance

Just Right

Low Bias
High Variance

Gradient Descent

Start with a cost function J(β):

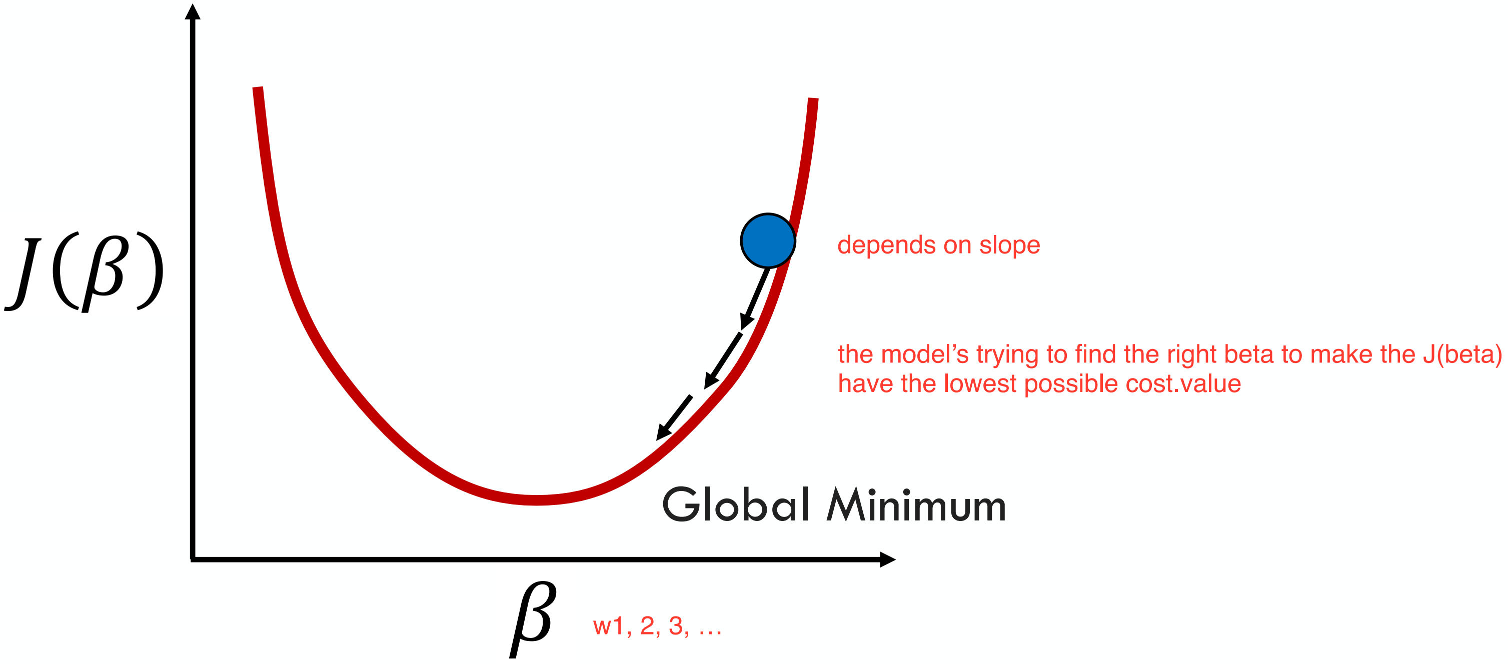

Then gradually move towards the minimum.

Gradient Descent with Linear Regression

- Now imagine there are two parameters (β_0, β_1)
- This is a more complicated surface on which the minimum must be found
- How can we do this without knowing what $J(\beta_0, \beta_1)$ looks like?

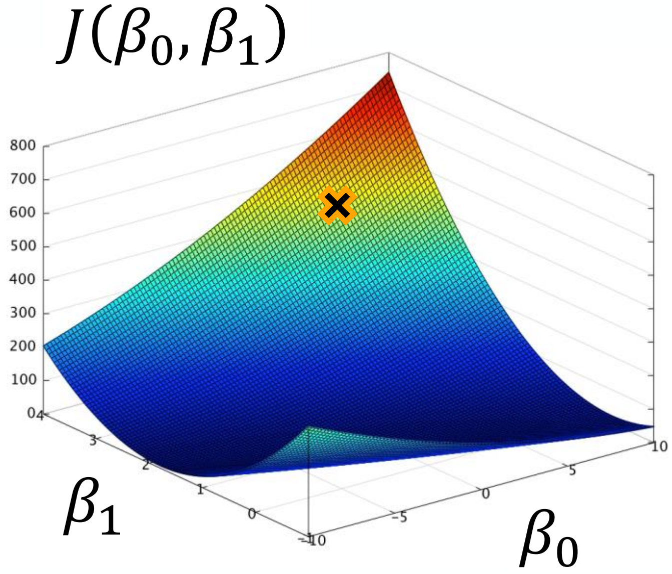

Gradient Descent with Linear Regression

- The gradient is a vector whose coordinates consist of the partial derivatives of the parameters

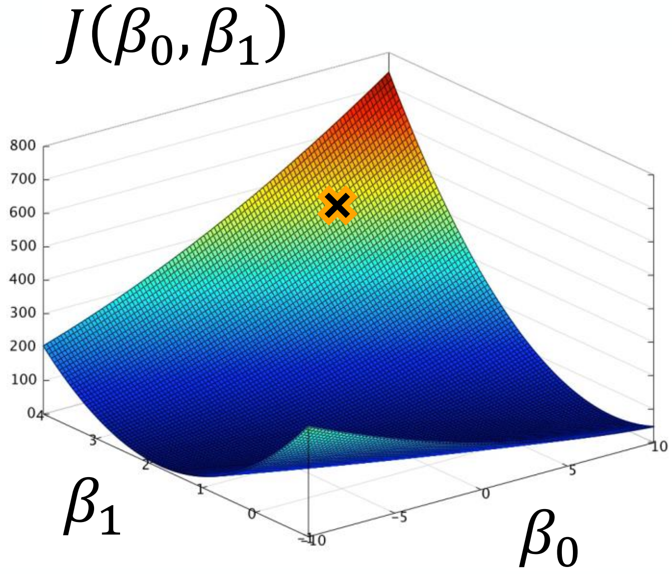

$$\nabla J(\beta_0, \dots, \beta_n) = < \frac{\partial J}{\partial \beta_0}, \dots, \frac{\partial J}{\partial \beta_n} >$$

Gradient Descent with Linear Regression

- Compute the gradient, $\nabla J(\beta_0, \beta_1)$, which points in the direction of the biggest increase!

Gradient (Slope)

- -$\nabla J(\beta_0, \beta_1)$ (negative gradient) points to the biggest decrease at that point!

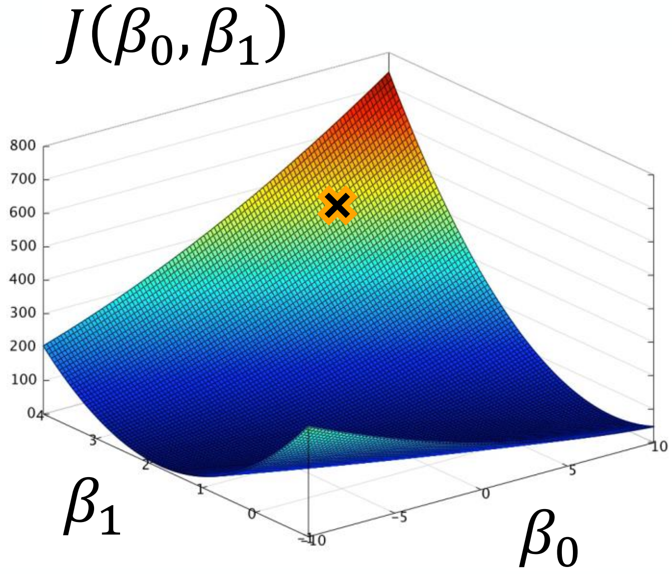

Gradient Descent with Linear Regression

- Then use the gradient (∇) and the cost function to calculate the next point (ω_1) from the current one (ω_0):

$$\omega_1 = \omega_0 - \alpha\nabla^{\text{Diff}}\frac{1}{2}\sum_{i=1}^{m}\left(\left(\beta_0 + \beta_1 x_{obs}^{(i)}\right) - y_{obs}^{(i)}\right)^2$$

- The learning rate (α) is a tunable parameter that determines step size

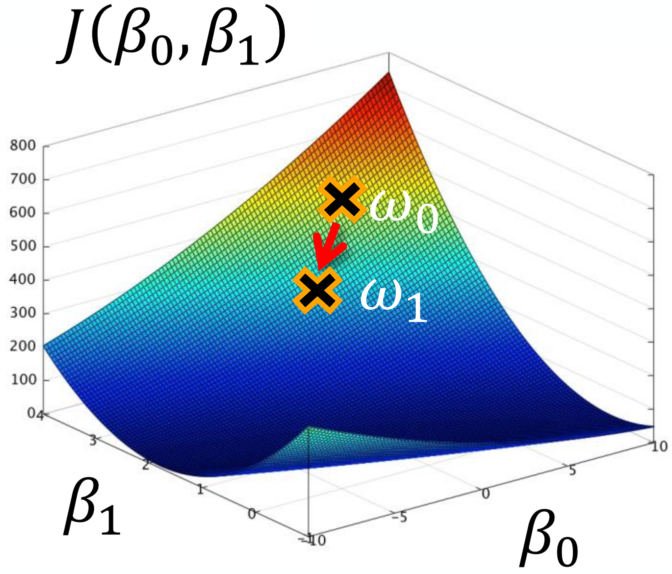

Gradient Descent with Linear Regression

- Each point can be iteratively calculated from the previous one

$$\omega_2 = \omega_1 - \alpha \nabla \frac{1}{2} \sum_{i=1}^{m} \left(\left(\beta_0 + \beta_1 x_{obs}^{(i)} \right) - y_{obs}^{(i)} \right)^2$$

$$\omega_3 = \omega_2 - \alpha \nabla \frac{1}{2} \sum_{i=1}^{m} \left(\left(\beta_0 + \beta_1 x_{obs}^{(i)} \right) - y_{obs}^{(i)} \right)^2$$

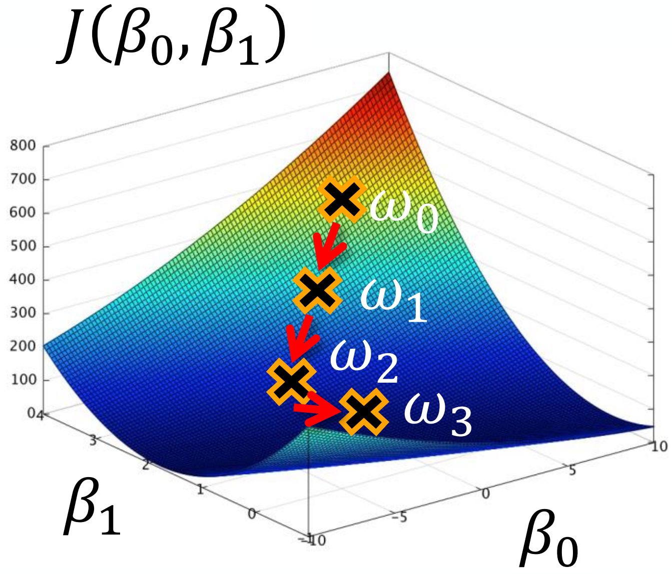

Stochastic Gradient Descent

Stochastic random only one data form all to calculate the gradient

- Use a single data point to determine the gradient and cost function instead of all the data

Full Batch

$$\omega_1 = \omega_0 - \alpha \nabla \frac{1}{2} \cancel{\sum_{i=1}^{m}} \left(\left(\beta_0 + \beta_1 x_{obs}^{(i)} \right) - y_{obs}^{(i)} \right)^2$$

$$\boxed{\omega_1 = \omega_0 - \alpha \nabla \frac{1}{2} \left(\left(\beta_0 + \beta_1 x_{obs}^{(0)} \right) - y_{obs}^{(0)} \right)^2}$$

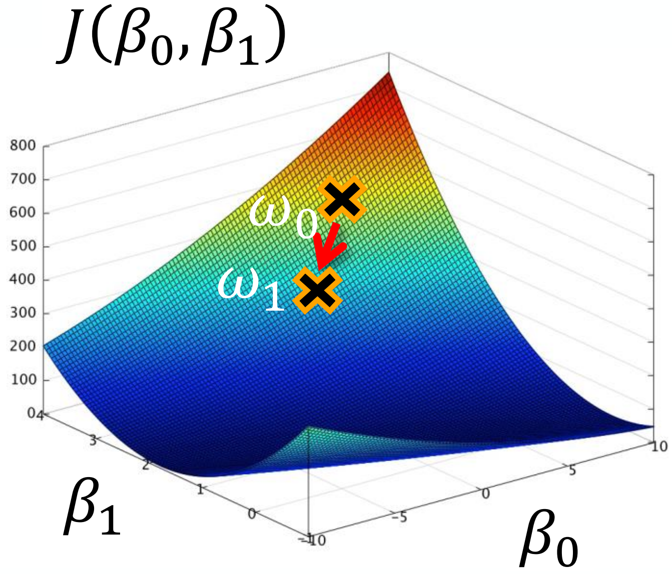

Stochastic Gradient Descent

- Use a single data point to determine the gradient and cost function instead of all the data

$$\omega_1 = \omega_0 - \alpha \nabla \frac{1}{2}\left(\left(\beta_0 + \beta_1 x_{obs}^{(0)}\right) - y_{obs}^{(0)}\right)^2$$

$$\dots$$

$$\omega_4 = \omega_3 - \alpha \nabla \frac{1}{2}\left(\left(\beta_0 + \beta_1 x_{obs}^{(3)}\right) - y_{obs}^{(3)}\right)^2$$

- Path is less direct due to noise in single data point—"stochastic"

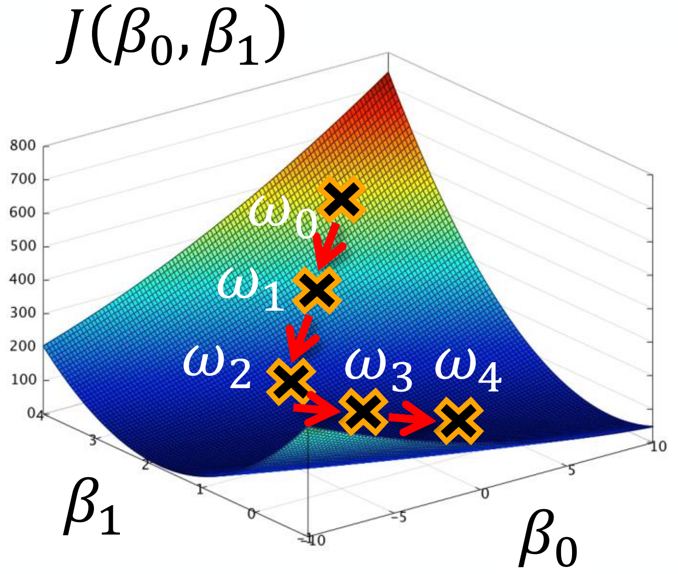

Mini Batch Gradient Descent

Something like 100 from 1,000

- Perform an update for every n training examples

$$\omega_1 = \omega_0 - \alpha\nabla\frac{1}{2}\sum_{i=1}^{n}\left(\left(\beta_0 + \beta_1 x_{obs}^{(i)}\right) - y_{obs}^{(i)}\right)^2$$

Best of both worlds:

- Reduced memory relative to "vanilla" gradient descent

- Less noisy than stochastic gradient descent

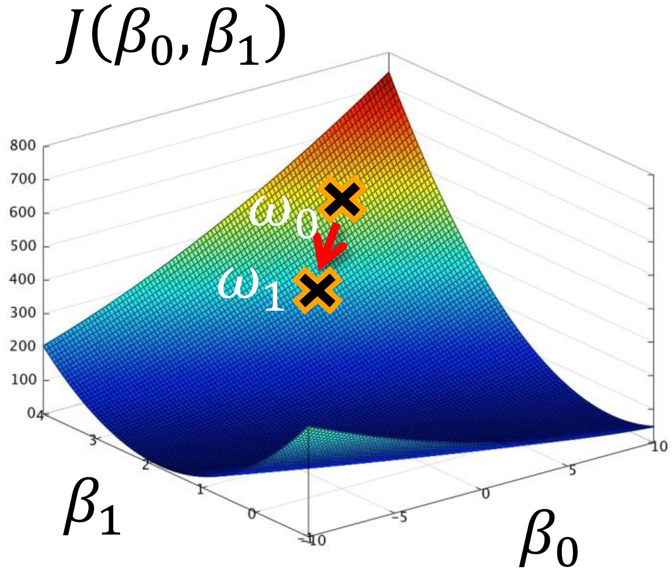

Software